

6.824 2015 Lecture 2: Infrastructure: RPC and threads

Remote Procedure Call (RPC)

a key piece of distrib sys machinery; all the labs use RPC

goal: easy-to-program network communication

hides most details of client/server communication

client call is much like ordinary procedure call

server handlers are much like ordinary procedures

RPC is widely used!

RPC ideally makes net communication look just like fn call:

Client:

z = fn(x, y)

Server:

fn(x, y) {

compute

return z

}

RPC aims for this level of transparency

Examples from lab 1:

DoJob

Register

RPC message diagram:

Client

Server

request--->

<---response

Software structure

client app

handlers

stubs

dispatcher

RPC lib

RPC lib

net ----- net

A few details:

Which server function (handler) to call?

Marshalling: format data into packets

Tricky for arrays, pointers, objects, &c

Go's RPC library is pretty powerful!

some things you cannot pass: e.g., channels, functions

Binding: how does client know who to talk to?

Maybe client supplies server host name

Maybe a name service maps service names to best server host

Threads:

Client often has many threads, so > 1 call outstanding, match up replies

Handlers may be slow, so server often runs each in a thread

RPC problem: what to do about failures?

e.g. lost packet, broken network, slow server, crashed server

What does a failure look like to the client RPC library?

Client never sees a response from the server

Client does **not** know if the server saw the request!

Maybe server/net failed just before sending reply

[diagram of lost reply]

Simplest scheme: "at least once" behavior

RPC library waits for response for a while

If none arrives, re-send the request

Do this a few times
 Still no response -- return an error to the application

Q: is "at least once" easy for applications to cope with?

Simple problem w/ at least once:
 client sends "deduct \$10 from bank account"

Q: what can go wrong with this client program?
 Put("k", 10) -- an RPC to set key's value in a DB server
 Put("k", 20) -- client then does a 2nd Put to same key
 [diagram, timeout, re-send, original arrives very late]

Q: is at-least-once ever OK?
 yes: if it's OK to repeat operations, e.g. read-only op
 yes: if application has its own plan for coping w/ duplicates
 which you will need for Lab 1

Better RPC behavior: "at most once"
 idea: server RPC code detects duplicate requests
 returns previous reply instead of re-running handler

Q: how to detect a duplicate request?
 client includes unique ID (XID) with each request
 uses same XID for re-send

server:
 if seen[xid]:
 r = old[xid]
 else
 r = handler()
 old[xid] = r
 seen[xid] = true

some at-most-once complexities
 this will come up in labs 2 and on
 how to ensure XID is unique?
 big random number?
 combine unique client ID (ip address?) with sequence #?
 server must eventually discard info about old RPCs
 when is discard safe?
 idea:
 unique client IDs
 per-client RPC sequence numbers
 client includes "seen all replies $\leq X$ " with every RPC
 much like TCP sequence #s and acks
 or only allow client one outstanding RPC at a time
 arrival of seq+1 allows server to discard all \leq seq
 or client agrees to keep retrying for < 5 minutes
 server discards after 5+ minutes
 how to handle dup req while original is still executing?
 server doesn't know reply yet; don't want to run twice
 idea: "pending" flag per executing RPC; wait or ignore

What if an at-most-once server crashes and re-starts?
 if at-most-once duplicate info in memory, server will forget
 and accept duplicate requests after re-start
 maybe it should write the duplicate info to disk?
 maybe replica server should also replicate duplicate info?

What about "exactly once"?
 at-most-once plus unbounded retries plus fault-tolerant service
 Lab 3

Go RPC is "at-most-once"

- open TCP connection

- write request to TCP connection

- TCP may retransmit, but server's TCP will filter out duplicates

- no retry in Go code (i.e. will NOT create 2nd TCP connection)

- Go RPC code returns an error if it doesn't get a reply

- perhaps after a timeout (from TCP)

- perhaps server didn't see request

- perhaps server processed request but server/net failed before reply came back

Go RPC's at-most-once isn't enough for Lab 1

- it only applies to a single RPC call

- if worker doesn't respond, the master re-send to it to another worker

- but original worker may have not failed, and is working on it too

- Go RPC can't detect this kind of duplicate

- No problem in lab 1, which handles at application level

- Lab 2 will explicitly detect duplicates

Threads

- threads are a fundamental server structuring tool

- you'll use them a lot in the labs

- they can be tricky

- useful with RPC

- Go calls them goroutines; everyone else calls them threads

Thread = "thread of control"

- threads allow one program to (logically) do many things at once

- the threads share memory

- each thread includes some per-thread state:

- program counter, registers, stack

Threading challenges:

- sharing data

- two threads modify the same variable at same time?

- one thread reads data that another thread is changing?

- these problems are often called races

- need to protect invariants on shared data

- use Go sync.Mutex

- coordination between threads

- e.g. wait for all Map threads to finish

- use Go channels

- deadlock

- thread 1 is waiting for thread 2

- thread 2 is waiting for thread 1

- easy detectable (unlike races)

- lock granularity

- coarse-grained -> simple, but little concurrency/parallelism

- fine-grained -> more concurrency, more races and deadlocks

- let's look at a toy RPC package to illustrate these problems

look at today's handout -- l-rpc.go

- it's a simplified RPC system

- illustrates threads, mutexes, channels

- it's a toy, though it does run

- assumes connection already open

- only supports an integer arg, integer reply

- omits error checks

```
struct ToyClient
```

```
    client RPC state
```

mutex per ToyClient
 connection to server (e.g. TCP socket)
 xid -- unique ID per call, to match reply to caller
 pending[] -- chan per thread waiting in Call()
 so client knows what to do with each arriving reply

Call

application calls `reply := client.Call(procNum, arg)`
 procNum indicates what function to run on server
 WriteRequest knows the format of an RPC msg
 basically just the arguments turned into bits in a packet
 Q: why the mutex in Call()? what does `mu.Lock()` do?
 Q: could we move "`xid := tc.xid`" outside the critical section?
 after all, we are not changing anything
 [diagram to illustrate]
 Q: do we need to WriteRequest inside the critical section?
 note: Go says you are responsible for preventing concurrent map ops
 that's one reason the update to pending is locked

Listener

runs as a background thread
 what is `<-` doing?
 not quite right that it may need to wait on chan for caller

Back to Call()...

Q: what if reply comes back very quickly?
 could Listener() see reply before `pending[xid]` entry exists?
 or before caller is waiting for channel?

Q: should we put `reply:=<-done` inside the critical section?
 why is it OK outside? after all, two threads use it.

Q: why mutex per ToyClient, rather than single mutex per whole RPC pkg?

Server's Dispatcher()

note that the Dispatcher echos the xid back to the client
 so that Listener knows which Call to wake up
 Q: why run the handler in a separate thread?
 Q: is it a problem that the dispatcher can reply out of order?

main()

note registering handler in `handlers[]`
 what will the program print?

Q: when to use channels vs shared memory + locks?
 here is my opinion
 use channels when you want one thread to explicitly wait for another
 often wait for a result, or wait for the next request
 e.g. when client Call() waits for Listener()
 use shared memory and locks when the threads are not intentionally
 directly interacting, but just happen to r/w the same data
 e.g. when Call() uses `tc.xid`
 but: they are fundamentally equivalent; either can always be used.

Go's "memory model" requires explicit synchronization to communicate!

This code is not correct:

```
var x int
done := false
go func() { x = f(...); done = true }
while done == false { }
```

it's very tempting to write, but the Go spec says it's undefined
use a channel or `sync.WaitGroup` instead

Study the Go tutorials on goroutines and channels