

6.824 2015 Lecture 3: Primary/Backup Replication

Today

- Replication
- Remus case study
- Lab 2 introduction

Fault tolerance

- we'd like a service that continues despite failures!
- available: still useable despite [some class of] failures
- correct: act just like a single server to clients
- very hard!
- very useful!

Need a failure model: what will we try to cope with?

- Independent fail-stop computer failure
 - Remus further assumes only one failure at a time
- Site-wide power failure (and eventual reboot)
- (Network partition)
- No bugs, no malice

Core idea: replication

- *Two* servers (or more)
- Each replica keeps state needed for the service
- If one replica fails, others can continue

Example: fault-tolerant MapReduce master

- lab 1 workers are already fault-tolerant, but not master
- master is a "single point of failure"
- can we have two masters, in case one fails?
- [diagram: M1, M2, workers]
- state:
 - worker list
 - which jobs done
 - which workers idle
 - TCP connection state
 - program counter

Big Questions:

- What state to replicate?
- How does replica get state?
- When to cut over to backup?
- Are anomalies visible at cut-over?
- How to repair / re-integrate?

Two main approaches:

- State transfer
 - "Primary" replica executes the service
 - Primary sends [new] state to backups
- Replicated state machine
 - All replicas execute all operations
 - If same start state,
 - same operations,
 - same order,
 - deterministic,
 - then same end state

State transfer is simpler

- But state may be large, slow to transfer
- Remus uses state transfer

Replicated state machine can be more efficient
 If operations are small compared to data
 But complex, e.g. order on multi-core, determinism
 Labs use replicated state machines

Remus: High Availability via Asynchronous Virtual Machine Replication
 NSDI 2008

Very ambitious system:

- Whole-system replication
- Completely transparent to applications and clients
- High availability for any existing software
- Would be magic if it worked well!
- Failure model:
 1. independent hardware faults
 2. site-wide power failure

Plan 1 (slow, broken):

[diagram: app, O/S, Remus underneath]
 two machines, primary and backup; plus net and other machines
 primary runs o/s and application s/w, talks to clients, &c
 backup does **not** initially execute o/s, applications, &c
 it only executes some Remus code
 a few times per second:
 pause primary
 copy entire RAM, registers, disk to backup
 resume primary
 if primary fails:
 start backup executing!

Q: is Plan 1 correct?

i.e. does it look just like a single reliable server?

Q: what will outside world see if primary fails and replica takes over?

will backup have same state as last visible on primary?
 might a client request be lost? executed twice?

Q: is Plan 1 efficient?

Can we eliminate the fact that backup **state** trails the primary?

Seems very hard!

Primary would have to tell backup (and wait) on every instruction.

Can we **conceal** the fact that backup's state lags primary?

Prevent outside world from **seeing** that backup is behind last primary state
 e.g. prevent primary sent RPC reply but backup state doesn't reflect that RPC
 e.g. MR Register RPC, which it would be bad for backup to forget

Idea: primary "holds" output until backup state catches up to output point
 e.g. primary receives RPC request, processes it, creates reply packet,
 but Remus holds reply packet until backup has received corresponding state update

Remus epochs, checkpoints

Clients: C1

	req1		reply1
Primary:	... E1 ...	pause E2	release pause
		ckpt	ok ckpt
Backup:	... (E0) ...	apply (E1)	

1. Primary runs for a while in Epoch 1, holding E1's output
2. Primary pauses
3. Primary sends RAM+disk changes to backup (in background)

4. Primary resumes execution in E2, holding E2's output
5. Backup copies all to separate RAM, then ACKs
6. Primary releases E1's output
7. Backup applies E1's changes to RAM and disk

If primary fails, backup finishes applying last epoch's disk+RAM, then starts executing

Q: any externally visible anomalies?

lost input/output?

repeated output?

Q: what if primary receives+executes a request, crashes before checkpoint?

backup won't have seen request!

Q: what if primary crashes after sending ckpt to backup,

but before releasing output?

Q: what if client doesn't use TCP -- doesn't re-transmit?

Q: what if primary fails while sending state to backup?

i.e. backup is mid-way through absorbing new state?

Q: are there situations in which Remus will incorrectly activate the backup?

i.e. primary is actually alive

network partition...

Q: when primary recovers, how does Remus restore replication?

needed, since eventually active ex-backup will itself fail

Q: what if *both* fail, e.g. site-wide power failure?

RAM content will be lost, but disks will probably survive

after power is restored, reboot guest from one of the disks

O/S and application recovery code will execute

disk must be "crash-consistent"

so probably not the backup disk if was in middle of installing checkpoint

disk shouldn't reflect any held outputs (... why not?)

so probably not the primary's disk if was executing

I do not understand this part of the paper (Section 2.5)

seems to be a window during which neither disk could be used if power failed

primary writes its disk during epoch

meanwhile backup applies last epoch's writes to its disk

Q: in what situations will Remus likely have good performance?

Q: in what situations will Remus likely have low performance?

Q: should epochs be short or long?

Remus Evaluation

summary: 1/2 to 1/4 native speed

checkpoints are big and take time to send

output hold limits speed at which clients can interact

Why so slow?

checkpoints are big and take time to generate and send

100ms for SPECweb2005 -- because many pages written

so inter-checkpoint intervals must be long

so output must be held for quite a while

so client interactions are slow

only 10 RPCs per second per client

How could one get better performance for replication?

- big savings possible with application-specific schemes:
 - just send state really needed by application, not all state
 - send state in optimized format, not whole pages
 - send operations if they are smaller than state
- likely **not** transparent to applications
 - and probably not to clients either

PRIMARY-BACKUP REPLICATION IN LAB 2

outline:

- simple key/value database
 - Get(k), Put(k, v), Append(k, v)
- primary and backup
 - replicate by primary sending each operation to backups
 - tolerate network problems, including partition
 - either keep going, correctly
 - or suspend operations until network is repaired
 - allow replacement of failed servers
 - you implement essentially all of this (unlike lab 1)

"view server" decides who p and b are

- main goal: avoid "split brain" — disagreement about who primary is
- clients and servers ask view server
- they don't make independent decisions

repair:

- view server can co-opt "idle" server as b after old b becomes p
- primary initializes new backup's state

key points:

1. only one primary at a time!
 2. the primary must have the latest state!
- we will work out some rules to ensure these

view server

- maintains a sequence of "views"

- view #, primary, backup

- 0: -- --

- 1: S1 --

- 2: S1 S2

- 4: S2 --

- 3: S2 S3

- monitors server liveness

- each server periodically sends a Ping RPC

- "dead" if missed N Pings in a row

- "live" after single Ping

- can be more than two servers Pinging view server

- if more than two, "idle" servers

- if primary is dead

- new view with previous backup as primary

- if backup is dead, or no backup

- new view with previously idle server as backup

- OK to have a view with just a primary, and no backup

- but -- if an idle server is available, make it the backup

how to ensure new primary has up-to-date replica of state?

- only promote previous backup

- i.e. don't make an idle server the primary

- backup must remember if it has been initialized by primary

if not, don't function as primary even if promoted!

Q: can more than one server think it is primary?

1: S1, S2

net broken, so viewserver thinks S1 dead but it's alive

2: S2, --

now S1 alive and not aware of view #2, so S1 still thinks it is primary

AND S2 alive and thinks it is primary

=> split brain, no good

how to ensure only one server acts as primary?

even though more than one may **think** it is primary

"acts as" == executes and responds to client requests

the basic idea:

1: S1 S2

2: S2 --

S1 still thinks it is primary

S1 must forward ops to S2

S2 thinks S2 is primary

so S2 must reject S1's forwarded ops

the rules:

1. primary in view i must have been primary or backup in view i-1
2. if you think you are primary, must wait for backup for each request
3. if you think you are not backup, reject forwarded requests
4. if you think you are not primary, reject direct client requests

so:

before S2 hears about view #2

S1 can process ops from clients, S2 will accept forwarded requests

S2 will reject ops from clients who have heard about view #2

after S2 hears about view #2

if S1 receives client request, it will forward, S2 will reject

so S1 can no longer act as primary

S1 will send error to client, client will ask vs for new view,

client will re-send to S2

the true moment of switch-over occurs when S2 hears about view #2

how can new backup get state?

e.g. all the keys and values

if S2 is backup in view i, but was not in view i-1,

S2 should ask primary to transfer the complete state

rule for state transfer:

every operation (Put/Get/Append) must be either before or after state xfer

== state xfer must be atomic w.r.t. operations

either

op is before, and xferred state reflects op

op is after, xferred state doesn't reflect op, prim forwards op after state

Q: does primary need to forward Get()s to backup?

after all, Get() doesn't change anything, so why does backup need to know?

and the extra RPC costs time

Q: how could we make primary-only Get()s work?

Q: are there cases when the lab 2 protocol cannot make forward progress?

View service fails

Primary fails before new backup gets state

We will start fixing those in lab 3