

6.824 2015 Lecture 17: PNUTS

Cooper et al, PNUTS: Yahoo!'s Hosted Data Serving Platform. VLDB 2008.

why this paper?

- same basic goals as Facebook/memcache and Spanner
 - multi-region data replication -- facing 100ms network delays
 - more principled design than Facebook/memcache, better semantics
 - simpler than Spanner
- apps control trade-off between consistency and performance

PNUTS' overall story similar to that of Spanner and Facebook/memcache

- data centers ("regions") all over the world
- each region has web servers and storage
- the big point: low latency via user contacting nearest region
- for small records: user profile, messages, shopping cart, friend list, &c

design overview

- [diagram: 3 regions, browsers, apps, tablet ctrlrs, routers, SUs, MBs]
- each region has all data
- each table partitioned by key over storage units
 - tablet servers + routers know the partition plan

why replicas of *all* data at multiple regions?

- any user can read anything quickly
 - good for data used by many users / many regions
 - cross-region reads would be slow (~100ms)
- failed SU repairable from another replica

drawbacks of a copy at each region?

- updates need to be sent to every region
- local reads will probably be stale
- updates from multiple regions need to be sorted out
 - keep replicas identical
 - don't lose updates (e.g. read-modify-write for counter)
- uses more disk space? not a problem for small records.

how do updates work?

- app server gets web request, needs to write data in PNUTS
- need to update every region!
- why not just have app logic send update to every region?
 - what if app crashes after updating only some regions?
 - what if concurrent updates to same record?

PNUTS has a "record master" for each record

- master imposes order on writes for each record
- responsible storage unit executes updates one at a time per record
- tells MB to broadcast updates in order to all regions

why per-record master?

- each record has a hidden column indicating region of record master
- finer-grained than Facebook/memcache master region
- hopefully get best of both worlds:
 - fast reads of local data
 - fast writes b/c master is often the local region

the complete update story (some guesses):

- app wants to update some columns of a record, knows key
 1. app sends update to local router
 2. router forwards to local SU1

3. SU1 looks up record master for key: region R2
4. SU1 sends update request to router in R2
5. R2 router forwards update to local SU2
6. SU2 sends update to local Message Broker (MB)
7. MB stores on disk + backup MB, sends vers # to original app
8. MB sends update to MB at every region
9. every region updates local copy (MB, router, SU)

MB is a neat idea

atomic: updates all replicas, or none

rather than app server updating replicas (crash...)

reliable: logs to disk, keeps trying, to cope with various failures

ordered: record's replicas eventually identical even w/ multiple writers

async: apps don't wait once write committed at MB

write order semantics

suppose each record in a table looks like:

Name Where What

Alice home sleeping

PNUTS preserves order of writes to a single record

e.g. Alice does

write(Alice.What, awake)

write(Alice.Where, work)

write RPC only returns after committed to record master's MB

readers may see home/sleeping or home/awake or work/awake

no-one will see work / sleeping

BUT writes to different records do not maintain order

e.g.

Write(Bob.What, on duty)

Write(Alice.What, off duty)

readers may see these in either order

i.e. no-one on duty; or two people on duty

if atomic hand-off important, need a single record with name of person

why is it OK for writes to take effect slowly?

fundamental benefit: reads are then very fast, since local

and reads usually greatly outnumber writes

PNUTS mitigates write cost:

application waits for MB commit but not propagation ("asynchronous")

master likely to be local (they claim 80% of the time)

so MB commit will often be quick

still, eval says write takes 300ms if master is remote!

down side: readers at non-master regions may see stale data

how stale might a non-master record be?

depends on how quickly MB sends updates to regions

guess: less than a second usually

longer if network slow/flaky, or MB busy

how do apps cope with potentially stale local replica?

sometimes stale is ok: looking at other people's profiles

sometimes stale is not ok: shopping cart after add/delete item

application gets to choose how consistent (section 2.2)

read-any(k)

read from local SU

fast but maybe stale

read-critical(k, required_version)

fast local read if local SU has vers \geq required_version

otherwise slow read from master SU

why: app knows it just wrote, wants value to reflect write

read-latest(k)

always read from master SU
 slow if master is remote!
 why: app needs fresh data

what if app needs to increment a counter stored in a record?
 is read-latest(k), increment, then write(k) enough?
 not if there might be concurrent updates!

test-and-set-write(version#, new value) gives you atomic update to one record
 master rejects the write if current version # != version#
 so if concurrent updates, one will lose and retry
 while(1):
 (x, ver) = read-latest(k)
 if(t-a-s-w(k, ver, x+1))
 break
 t-a-s-w is fully general for single-record atomic read-modify-write

why not atomic multi-record writes?
 e.g. for bank transfers: Alice -= \$10 ; Bob += \$10
 would that be easy to add to PNUTS?

The Question

how does PNUTS cope with Example 1 (page 2)
 Initially Alice's mother is in Alice's ACL, so mother can see photos
 1. Alice removes her mother from ACL
 2. Alice posts spring-break photos
 could her mother see update #2 but not update #1?
 ACL and photo list must be in the same record
 since PNUTS guarantees order only for updates to same record
 How could a storage system get this wrong?
 No ordering through single master (e.g. Dynamo)

What if Alice's mother does two reads:

1. reads the ACL
 2. reads the photo list
- #1 could yield old ACL (with mother in it); #2 could yield new photo
 so: mother must read Alice's record just once

how to change record's master if no failures?

e.g. I move from Boston to LA
 perhaps just update the record, via old master?
 since ID of master region is stored in the record
 old master announces change over MB
 a few subsequent updates might go to the old master SU
 it will reject them, app retries and finds new master?

what about tolerating failures?

app server crashes midway through a set of updates
 not a transaction, so only some of writes will happen
 but master SU/MB either did or didn't get each write
 so each write happens at all regions, or none

SU down briefly, or network temporarily broken/lossy
 (I'm guessing here, could be wrong)
 MB keeps trying until SU acks
 SU shouldn't ACK until safely on disk

SU loses disk contents, or doesn't automatically reboot
 need to restore disk content from SUs at other regions
 1. subscribe to MB feed, and save them for now

2. copy content from SU at another region
3. replay saved MB updates

puzzle:

how to ensure we didn't miss any MB updates for this SU?
 e.g. subscribe to MB at time=100, but source SU only saw through 90?
 will replay apply updates twice? is that harmful?
 paper mentions sending checkpoint message through MB
 maybe fetch copy as of when the checkpoint arrived
 and only replay after the checkpoint
 BUT no ordering among MB streams from multiple regions

SU crashes in the middle of an update

does SU update local disk, then send to MB?
 or does SU forward to MB, and only apply to disk after MB commits?

MB crashes after accepting update

logs to disks on two MB servers before ACKing
 MB recovery looks at log, (re)sends logged msgs
 record master SU maybe re-sends an update if MB crash before ACK
 maybe record version #s will allow SUs to ignore duplicate

record's master region loses network connection

can other regions designate a replacement RM?
 no: original RM's MB may have logged updates, only some sent out
 do other regions have to wait indefinitely? yes
 this is one price of ordered updates / strict-ish consistency

now performance

evaluation focuses on latency, not throughput

why are they apparently worried about latency but not throughput?
 maybe throughput can be indefinitely increased by adding more SUs
 whereas latency is harder to reduce

big performance question: why isn't the MB a terrible bottleneck?

all writes funnel through a single pair of MBs
 which log every write to disk
 i do not know how they avoid this MB bottleneck

5.2: time for an insert while busy

probably measuring time until client sees commit reply from MB
 depends on how far away Record Master is
 RM local: 75.6 ms
 RM nearby: 131.5 ms
 RM other coast: 315.5 ms

Why 75 ms for write to local RM?

speed-of-light delay?
 no: local
 queuing, waiting for other client's operations?
 no: they imply 100 clients was max that didn't cause delay to rise
 End of 5.2 suggests 40 ms of 75 ms in in SU
 how could it take 40 ms?
 each key/value is one file?
 creating a file takes 3 disk writes (directory, inode, content)?
 what's the other 35 ms?
 MB disk write?

5.3 / Figure 3: effect of increasing request rate

what do we expect for graph w/ x-axis req rate, y-axis latency?
 system has some inherent capacity, e.g. total disk seeks/second

for lower rates, constant latency
for higher rates, queue grows rapidly, avg latency blows up
blow-up should be near max capacity of h/w
e.g. # disk arms / seek time
we don't see a blow-up in Figure 3
end of 5.3 says clients too slow
text says max possible rate was about 3000/second
10% writes, so 300 writes/second
5 SU per region, so 60 writes/SU/second
about right if each write does a random disk I/O
but you'll need lots of SUs for millions of active users
mystery: how is MB able to log 300 writes/second?
maybe each region's MB only logs 1/3 of writes, i.e. 100 second?

stepping back, what were PNUTS key design decisions?

1. replication of all data at multiple regions
fast reads, slow but async writes
2. relaxed consistency -- stale reads
needed if you want fast reads and async writes
3. sequence all writes thru record master
good for consistency; bad for latency
4. only single-row write semantics (order and t-a-s-w mini-transactions)
simplifies system; maybe awkward for applications

Next: Dynamo, a very different design

no master -- any region can update any value at any time
eventual consistency
tree of versions if network partitions
readers must reconcile versions