

6.824 2015 Lecture 21: Distributed Optimistic Concurrency Control

Paper: Efficient Optimistic Concurrency Control using Loosely Synchronized Clocks, by Adya, Gruber, Liskov and Maheshwari.

Why this paper?

- to look at optimistic concurrency control (OCC)
- OCC might help us get large scale, high speed, *and* good semantics

Thor overview

- [clients, client caches, servers A-M N-Z]
- data sharded over servers
- code runs in clients (not like Argus; not an RPC system)
- clients read/write DB records from servers
- clients cache data locally for fast access
- on client cache miss, fetch from server

Thor arrangement is fairly close to modern big web site habits

- clients, local fast cache, slower DB servers
- like Facebook/memcache paper
- but Thor has much better semantics

Thor programs use fully general transactions

- multi-operation
- serializable
- so can do bank xfers w/o losing money, &c

Client caching makes transactions tricky

- writes have to invalidate cached copies
- how to cope with reads of stale cached data?
- how to cope with read-modify-write races?
- clients could lock before using each record
- but that's slow -- probably need to contact server
- wrecks the whole point of fast local caching in clients
- (though caching read locks might be OK, as in paper Eval)

Thor uses optimistic concurrency control (OCC)

- an idea from the early 1980s
- just read and write the local copy
- don't worry about other transactions until commit
- when transaction wants to commit:
 - send read/write info to server for "validation"
 - validation decides if OK to commit -- if serializable
 - if yes, send invalidates to clients with cached copies of written records
 - if no, abort, discard writes
- optimistic b/c hopes for no conflict
- if turns out to be true, fast!
- if false, validation can detect, but slow

What should validation do?

- it looks at what the executing transactions read and wrote
- decides if there's a serial execution order that would have gotten the same results as the actual concurrent execution
- there are many OCC validation algorithms!
- i will outline a few, leading up to Thor's

Validation scheme #1

- a single validation server
- clients tell validation server the read and write VALUES seen by each transaction that wants to commit

"read set" and "write set"
 validation must decide:
 would the results be serializable if we let these
 transactions commit?
 scheme #1 shuffles the transactions, looking for a serial order
 in which each read sees the value written by the most
 recent write; if one exists, the execution was serializable.

Validation example 1:

initially, x=0 y=0 z=0
 T1: Rx0 Wx1
 T2: Rz0 Wz9
 T3: Ry1 Rx1
 T4: Rx0 Wyl
 validation needs to decide if this execution (reads, writes)
 is equivalent to some serial order
 yes: one such order is T4, T1, T3, T2; says yes to all
 (really T2 can go anywhere)
 note this scheme is far more permissive than Thor's
 e.g. it allows transactions to see uncommitted writes

OCC is neat b/c transactions didn't need to lock!
 so they can run quickly from client caches
 just one msg exchange w/ validator per transaction
 not one locking exchange per record used
 OCC excellent for T2 which didn't conflict with anything
 we got lucky for T1 T3 T4, which do conflict

Validation example 2 -- sometimes must abort:

initially, x=0 y=0
 T1: Rx0 Wx1
 T2: Rx0 Wyl
 T3: Ry0 Rx1
 values not consistent w/ any serial order!
 T1 -> T3 (via x)
 T3 -> T2 (via y)
 T2 -> T1 (via x)
 there's a cycle, so not the same as any serial execution
 perhaps T3 read a stale y=0 from cache
 or T2 read a stale x=0 from cache
 in this case validation can abort one of them
 then others are OK to commit
 e.g. abort T2
 then T1, T3 is OK (but not T3, T1)

How should client handle abort?

roll back the program (including writes to program variables)
 re-run from start of transaction
 hopefully won't be conflicts the second time
 OCC is best when conflicts are uncommon!

Do we need to validate read-only transactions?

example:
 initially x=0 y=0
 T1: Wx1
 T2: Rx1 Wy2
 T3: Ry2 Rx0
 i.e. T3 read a stale x=0 from its cache, hadn't yet seen invalidate.
 need to validate in order to abort T3.
 other OCC schemes can avoid validating read-only transactions
 keep multiple versions -- but Thor and my schemes don't

Is OCC better than locking?

yes, if few conflicts

avoids lock msgs, clients don't have to wait for locks

no, if many conflicts

OCC aborts, must re-start, perhaps many times

locking waits

example: simultaneous increment

locking:

T1: Rx0 Wx1

T2: -----Rx1 Wx2

OCC:

T1: Rx0 Wx1

T2: Rx0 Wx1

fast but wrong; must abort one

We really want **distributed** OCC validation

split storage and validation load over servers

each storage server sees only xactions that use its data

each storage server validates just its part of the xaction

two-phase commit (2PC) to check that they all say "yes"

only really commit if all relevant servers say "yes"

Can we just distribute validation scheme #1?

imagine server S1 knows about x, server S2 knows about y

example 2 again

T1: Rx0 Wx1

T2: Rx0 Wy1

T3: Ry0 Rx1

S1 validates just x information:

T1: Rx0 Wx1

T2: Rx0

T3: Rx1

answer is "yes" (T2 T1 T3)

S2 validates just y information:

T2: Wy1

T3: Ry0

answer is "yes" (T3 T2)

but we know the real answer is "no"

So simple distributed validation does not work

the validators must choose consistent orders!

Validation scheme #2

Idea: client (or TC) chooses timestamp for committing xaction

from loosely synchronized clocks, as in Thor

validation checks that reads and writes are consistent with TS order

solves distrib validation problem:

timestamps tell the validators the order to check

so "yes" votes will refer to the same order

Example 2 again, with timestamps:

T1@100: Rx0 Wx1

T2@110: Rx0 Wy1

T3@105: Ry0 Rx1

S1 validates just x information:

T1@100: Rx0 Wx1

T2@110: Rx0

T3@105: Rx1

timestamps say order must be T1, T3, T2

does not validate! T2 should have seen x=1

S2 validates just y information:

T2@110: Wy1

T3@105: Ry0

timestamps say order must be T3, T2

validates!

S1 says no, S2 says yes, two-phase commit coordinator will abort

What have we given up by requiring timestamp order?

example:

T1@100: Rx0 Wx1

T2@50: Rx1 Wx2

T2 follows T1 in real time, and sees T1's write

but T2 will abort, since TS says T2 comes first, so T1 should have seen x=2

could have committed, since T1 then T2 works

this will happen if client clocks are too far off

if T1's client clock is ahead, or T2's behind

so: requiring TS order can abort unnecessarily

b/c validation no longer **searching** for an order that works

instead merely **checking** that TS order consistent w/ reads, writes

we've given up some optimism by requiring TS order

maybe not a problem if clocks closely synched

maybe not a problem if conflicts are rare

Problem with schemes so far:

commit messages contained **values**, which can be big

could instead use version numbers to check whether

later xaction read earlier xaction's write

let's use writing xaction's TS as record version number

Validation scheme #4

tag each DB record (and cached record) with TS of xaction that last wrote it

validation requests carry TS of each record read

Our example for scheme #4:

all values start with timestamp 0

T1@100: Rx@0 Wx

T2@110: Rx@0 Wy

T3@105: Ry@0 Rx@100

note:

reads have timestamp that was in read record, not value

writes don't include either value or timestamp

S1 validates just x information:

orders the transactions by timestamp:

T1@100: Rx@0 Wx

T3@105: Rx@100

T2@110: Rx@0

the question: does each read see the most recent write?

T3 is ok, but T2 is not

S2 validates just y information:

again, sort by TS, check each read saw latest write:

T3@105: Ry@0

T2@110: Wy

this does validate

so scheme #4 abort, correctly, reasoning only about version TSs

what have we give up by thinking about version #s rather than values?

maybe version numbers are different but values are the same

e. g.

T1@100: Wx1

T2@110: Wx2

T3@120: Wx1

T4@130: Rx1@100
 timestamps say we should abort T4 b/c read a stale version
 should have read T3's write
 so scheme #4 will abort
 but T4 read the correct value -- x=1
 so abort wasn't necessary

Problem: per-record timestamp might use too much storage space
 Thor wants to avoid space overhead
 maybe important, maybe not

Validation scheme #5

Thor's invalidation scheme: no timestamps on records
 how can validation detect that a transaction read stale data?
 it read stale data b/c earlier xaction's invalidation hadn't yet arrived!
 so server can track invalidation msgs that might not have arrived yet
 "invalid set" -- one per client
 delete invalid set entry when client ACKs invalidation msg
 server aborts committing xaction if it read record in client's invalid set
 client aborts running xaction if it read record mentioned in invalidation

Example use of invalid set

[timeline]

Client C1:

T2@105 ... Rx ... 2PC commit point
 imagine that client acts as 2PC coordinator

Server:

VQ: T1@100 Wx
 T1 committed, x in C1's invalid set
 server has sent invalidation message to C1

Three cases:

1. invalidation arrives before T2 reads
 Rx will miss in client cache, read from data from server
 client will (probably) return ACK before T2 commits
 server won't abort T2
2. invalidation arrives after T2 reads, before commit point
 client will abort T2 in response to invalidation
3. invalidation arrives after 2PC commit point
 i.e. after all servers replied to prepare
 this means the client was still in the invalid set when
 the server tried to validate the transaction
 so the server aborted, so the client will abort too
 so: Thor's validation detects stale reads w/o timestamp on each record

Performance

Look at Figure 5

AOCC is Thor

comparing to ACBL: client talks to srvr to get write-locks,
 and to commit non-r/o xactions, but can cache read locks along with data
 why does Thor (AOCC) have higher throughput?

fewer msgs; commit only, no lock msgs

why does Thor throughput go up for a while w/ more clients?

apparently a single client can't keep all resources busy
 maybe due to network RTT?

maybe due to client processing time? or think time?

more clients -> more parallel xactions -> more completed

why does Thor throughput level off?

maybe 15 clients is enough to saturate server disk or CPU
 abt 100 xactions/second, about right for writing disk

why does Thor throughput *drop* with many clients?
more clients means more concurrent xactions at any given time
more concurrency means more chance of conflict
for OCC, more conflict means more aborts, so more wasted CPU

Conclusions

fast client caching + transactions would be excellent
distributed OCC very interesting, still an open research area
avoiding per-record version #s doesn't seem compelling
Thor's use of time was influential, e.g. Spanner