

6.824 2015 Lecture 23: Bitcoin

Bitcoin: A Peer-to-Peer Electronic Cash System, by Satoshi Nakamoto
2008

why aren't credit cards the perfect digital money?

- + work online
- + hard to steal (a complex situation)
- + can cancel transactions, call customer support
- relies on 3rd parties to verify (banks)
b/c users cannot directly verify charges
- 3% fee
- long settling time
- hard to become a merchant
- + tied to currency controlled by government

bitcoin: e-money without a central trusted party
a public ledger: anyone can verify transactions

what's hard technically?

- forgery
- double spending
- theft

what's hard socially/economically?

- why does it have value?
- how to pay for infrastructure?
- monetary policy (intentional inflation &c)
- laws (taxes, laundering, drugs, terrorists)

let's design OneBit, a simple e-money system

- to illustrate a public, verifiable ledger using transaction chain
- each user owns some coins
- single server -- OneBank -- everyone talks to it
- OneBank records all transactions

OneBit transactions

- every coin has a chain of transaction records
 - one for each time this coin was transferred as payment
- OneBank maintains the complete chain for each coin
- chain helps ensure that only the current owner spends

what's in a OneBit transaction record?

- public key of new owner
- hash of this coin's previous transaction record
- signed by private key of previous owner
- (Bitcoin is much more complex: amount (fractional), multiple in/out, ...)

OneBit example:

- Y owns a coin, previously given to it by X:
T7: pub(Y), hash(T6), sig(X)
- Y buys a hamburger from Z and pays with this coin
- Z tells Y Z's public key ("address")
- Y creates a new transaction and signs it
T8: pub(Z), hash(T7), sig(Y)
- OneBank verifies that:
 - no other transaction mentions hash(T7),
 - T8's sig() corresponds to T7's pub()
- Z waits until OneBank has seen/verified T8,
verifies that T8's pub() is Z's public key,

then Z gives hamburger to Y

why include $\text{pub}(Z)$?

why sign with $\text{sig}(Y)$?

why include $\text{hash}(T)$?

$\text{hash}(T)$ identifies the exact coin to be spent

a coin ID scheme might be ambiguous if Y owned this coin previously

where is Z's resulting coin value "stored"?

coin balance = unspent xaction

the "identity" of a coin is the (hash of) its most recent xaction

Z "owns" the coin: has private key that allows Z to make next xaction

does OneBit's transaction chain prevent stealing?

current owner's private key needed to sign next transaction

danger: attacker can steal Z's private key

Z uses private key a lot, so probably on his PC, easy to steal?

a significant problem for BitCoin in practice

what if OneBank is corrupt?

it can't forge transactions (doesn't know private keys)

but it can help people double-spend!

double-spending with OneBit

suppose OneBank is cooperating with Y to cheat Z or Q

Y creates two transactions for same coin: $Y \rightarrow Z$, $Y \rightarrow Q$

both have has pointing to same current end of chain

OneBank shows chain ending in $Y \rightarrow Z$ to Z, and $Y \rightarrow Q$ to Q

both transactions look good, including signatures and hash

now both Z and Q will give hamburgers to Y

why was double-spending possible?

OneBank can *hide* some information,

or selectively reveal it

what's needed?

many servers ("peers")

send all transactions to all peers

much harder for a few bad peers to hide transactions

conventions to un-fork if problems do arise

the BitCoin block chain

single block chain contains transactions on all coins

a copy stored in each peer

so each peer can validate new transactions against old ones

each block:

$\text{hash}(\text{prevblock})$

set of transactions

"nonce" (not quite a nonce in the usual cryptographic sense)

current time (wall clock timestamp)

a transaction isn't real unless it's in the block chain

new block every 10 minutes containing xactions since prev block

who creates each new block?

all the peers try

requirement: $\text{hash}(\text{block}) < \text{"target"}$

peers try nonce values until this works out

can't predict a winning nonce, since cryptographic hash

trying one nonce is fast, but most nonces won't work

it would take one CPU months to create one block
 but thousands of peers are working on it
 such that expected time to first to find is about 10 minutes
 the winner sends the new block to all peers
 (this is part of "mining")

how do transactions work w/ block chain?

start: all peers know ...<-B5
 and are working on B6 (trying different nonces)
 Z sends public key (address) to Y
 Y sends Y->Z transaction to peers, which flood it
 peers buffer the transaction until B6 computed
 peers that heard Y->Z include it in next block
 so eventually ...<-B5<-B6<-B7, where B7 includes Y->Z

what if bad person wants to double-spend?

start with block chain ...<-B6
 Y gets Y->Z into block chain
 ...<-B6<-BZ (BZ contains Y->Z)
 Z will see Y->Z in chain and give Y a hamburger
 can Y create ...<-B6<-BQ
 and persuade peers to accept it in place of ...<-B6<-BZ?

when will a peer accept chain CX it hears from another peer?

suppose peer already knows of chain CY
 it will accept CX if $\text{len}(\text{CX}) > \text{len}(\text{CY})$
 and if CX passes some validity tests
 will not accept if $\text{len}(\text{CX}) = \text{len}(\text{CY})$: first chain of same length wins

so attacker needs a longer chain to double-spend

e.g. ...<-B6<-BQ<-B8, which is longer than ...<-B6<-BZ
 and must create it in less than 10 minutes
 before main chain grows by another block
 10 minutes is the time it takes the 1000s of honest peers
 to find one block
 if the attacker has just one CPU, will take months to create the blocks
 by that time the main chain will be much longer
 no peer will accept the attacker's shorter chain
 attacker has no chance
 if the attacker has 1000s of CPUs -- more than all the honest
 bitcoin peers -- then the attacker can double spend

summary:

attacker can force honest peers to switch chains if attacker
 controls majority of peer CPU power

how long should Z wait before giving Y the hamburger?

until Z sees Y flood the transaction to many peers?
 no -- not in the chain, Y might flood conflicting xaction
 maybe OK for low-value transactions
 until Z sees chain ...<-BZ?
 maybe
 risky -- non-zero chance that some other chain will win
 i.e. some lucky machine discovered a few blocks in a row
 quickly, but its network msgs have so far been lost
 perhaps that chain won't have Y->Z
 probability goes down rapidly with number of blocks
 until Z sees chain with multiple blocks after BZ?
 yes -- slim chance attacker with few CPUs could catch up

validation checks:

much of burden is on (honest) peers, to check new xactions/blocks
to avoid ever having to scan the whole block chain
and so that clients don't have to maintain the whole block chain
peer, new xaction:
no other transaction refers to same previous transaction
signature is by private key of previous transaction
[then will add transaction to txn list for new block being mined]
peer, new block:
hash value < target (i.e. nonce is right, proof of work)
previous block hash exists
new chain longer than current chain
all transactions in block are valid
Z:
Y->Z is in a recent block
Z's public key / address is in the transaction
multiple peers have accepted that block
there's several more blocks in the chain
(other stuff has to be checked as well, lots of details)

where does each bitcoin originally come from?
each time a peer creates a block, it gets 25 bitcoins
assuming it is the winner for that block
it puts its public key in a special transaction in the block
this is incentive for people to operate bitcoin peers
but that number halves every 210,000 blocks (abt 4 years)
the point: motivate people to run peers

Q: how do peers communicate / find each other?

Q: what prevents a bad peer from modifying an existing block?

Q: what if two nodes disagree on the validity of a transaction?
(slight implementation differences between software versions)

Q: 10 minutes is annoying; could it be made much shorter?

Q: are transactions anonymous? pseudonymous? analogy: IP addresses.

Q: can bitcoins be stolen?

Q: if I steal bitcoins, is it safe to spend them?

Q: what can adversary do with a majority of CPU power in the world?
can double-spend
cannot steal others' bitcoins
can prevent xaction from entering chain
can revert past xactions (by building longer chain from before that block)

Q: how rich are you likely to get with one machine mining?

Q: if more people (CPUs) mine, will that create new bitcoin faster?
important use of block timestamps: control difficulty (hash target)

Q: why mine at all? why not start with a fixed number of coins?

Q: why does it make sense for the mining reward to decrease with time?

Q: is it a problem that there will be a fixed number of coins?

Q: what if the real economy grows (or shrinks)?

Q: why do bitcoins have value?

e.g., 1 BTC appears to be around US\$242 on may 12th 2015.

Q: will we still need banks, credit cards, &c?

today, dollar bills are only a small fraction of total money

same may be true of bitcoin

so properties of bitcoin (anonymity &c) may not be very important

Q: what are the benefits of banks, credit cards, &c?

disputes (Z takes payment and does not give hamburger to Y)

loss / recovery (user cannot find their private key)

Q: will bitcoin scale well?

as transaction rate increases?

claim CPU limits to 4,000 tps (signature checks)

more than Visa but less than cash

as block chain length increases?

do you ever need to look at very old blocks?

do you ever need to xfer the whole block chain?

merkle tree: block headers vs txn data.

key idea: block chain

public ledger is a great idea

decentralization might be good

mining seems imperfect, but does avoid centralized trust

tying ledger to new currency seems bad